


The third front end over the flint bucket format, split out of the flint-fuse architecture doc: if the workspace fits on disk, no interception is needed at all. Materialize plain files at pod start, publish changed files at cadence with an unprivileged sidecar, and put write authority in a stateless gateway. Most of FUSE's value at a fraction of its risk surface — same bucket format, same engine.




This is a third front end, not a FUSE configuration — the app touches plain files on a real local filesystem, and the sidecar is an ordinary unprivileged program running the same engine verbs: `import` at start (bytes instead of stubs), the flush barrier's whole-file lane at cadence (driven by an mtime/size scan against the manifest instead of capture interception), the same epoch, claim, and manifest machinery. For the agent-workspace shape — code, configs, artifacts, GiB-scale — it is most of FUSE's value with none of FUSE's hard parts: the entire privileged/mount-propagation/ENOTCONN/mount-broker risk section of the FUSE architecture stops applying, and the kernel floor drops to nothing. **Two corrections from the drills.** This page originally read “gateway down ⇒ publishes pause”; chaos leg C8 measured the opposite — barriers continue, because the shipped sidecar writes its manifest, window and epoch cells straight to the store. All four stated outage effects belong to the **proxy**; the gateway costs only the fourth. And **grade 2 is the plan of record, not yet a build**: the S3 proxy IS the intended deployment — a project-scoped door holding the real credentials, so no pod stores one at all — but nothing in the tree mints a presigned URL today, and the interim path gives the sidecar a scoped credential of its own. Read every “pods hold zero credentials” claim on these pages as the target the proxy delivers; today the agent container holds none, which is the half that already matters for blast radius. **The gateway restores the security posture the review said direct mode gives up.** With grade 2, “the bucket trusts one principal” is true again — pods hold no write credentials, the cooperative CAS regime becomes enforced arbitration (a non-holder is refused a URL, not merely detected after the fact), revocation returns to seconds, and per-share metering exists again. Its failure mode is publish-pause with growing RPO while pods keep serving — the softest failure shape any component on these pages has. **One dialect everywhere:** this also argues for converging the hub's file API onto the S3 REST surface (the libflint plan's §8 sidecar adapter, built once, placed three ways — hub listener, localhost sidecar, gateway), so the live view and the published view speak the same protocol and the same tools work against both. **Recommendation:** lean is the direct-mode entry point; FUSE is the escalation for volumes too large or too lazy to materialize; the hub is unchanged. Same bucket, three front ends, one engine. Implementation plan: docs/plans/flint-lean-plan.md.

The decision of record: every REST surface in the system speaks the S3 dialect — the hub's file API converges onto it, the gateway serves it, and direct-mode non-mount access (lean and FUSE alike) is plain S3 against the bucket. One library, three placements; the consistency contracts say which view a client is reading.



two views, one protocol — the consistency contracts tell them apart



direct-mode REST access — lean AND FUSE variants alike — is the bucket endpoint: plain S3, no flint server anywhere in the read path. the same rclone / boto3 / aws-cli invocation works against either endpoint; only the URL and the consistency contract change.

Why S3 and not a nicer bespoke API: the ecosystem is the feature — every language has a mature S3 client, presigned flows are understood by every CI system and browser uploader, and the tools people already point at buckets (rclone, boto3, cyberduck, Velero) work against a live flint share the day the dialect ships. The file API was already shaped on S3 semantics — whole-file GET/PUT with ETag and If-Match (the v1.30 conditional-writes work maps one-to-one onto S3's own conditional headers) — so this is a dialect change, not an architecture change. **The one thing the hub's live view offers that no real S3 can:** assembly under a temp name completed by an atomic RENAME under a deny-both OPEN — an atomic, exclusive multipart-complete, the differentiator the libflint plan already claims for the localhost gateway, inherited here. **Migration is dual-serve:** the hub answers both /files and the S3 dialect for one release; the front door moves; /files deprecates. Bearer auth stays alongside SigV4 so existing gateway tokens keep working, and revocation keeps its seconds-scale semantics. The activity classification is the one place the dialect must be flint-aware from day one — the idle ladder's lesson stands: a conditional GET answering 304 pins a share exactly as hard as a 200, so what counts as "activity" is declared per verb, not inferred.

Pages 1–2 describe publishing **at cadence**: an mtime/size scan on a timer. That is still the default, and it is still the whole story for a workspace nobody is watching. What it could not express is a coherent point — “these files, together, now” — so an agent finishing a unit of work had no way to say so and no way to learn when its bytes were durable. The boundary verbs are that sentence, written as files in the workspace the agent already has. Nothing new is mounted, no port is opened, and a workspace that ignores them pays nothing: measured at **20 bucket requests per 22 s idle with the verbs off and 20 with them on**.



Lean is not a general filesystem and does not try to be. It is for **one agent, one workspace, one writer, a workspace that fits on disk** — the harness shape: clone-ish tree of code and artifacts, an agent that reads and writes it locally at full speed, and a durable point somebody else can read when the work is done. Everything below is the boundary of that claim, in measured numbers, so the answer to “will this work for us” is arithmetic rather than opinion.

three questions, in order

1 · does the workspace fit the disk?

the whole tree is materialized — no stubs, no eviction.
No ⇒ FUSE (its own doc). This is the hard gate.

2 · is one writer per subtree enough?

a lease admits exactly one sidecar; a second is refused, not merged. No ⇒ the hub, for live shared POSIX.

3 · is snapshot freshness acceptable?

readers see the last published boundary, not your last write. No ⇒ the hub. Yes ⇒ lean, and read on.

the loop — one agent, start to durable

pod starts

webhook injects the sidecar;
checkout runs BEFORE the
agent's first line

agent works

plain local files, full POSIX,
zero interception — this is
the part that is just fast

agent declares

echo > .flint/publish
one file write — the whole
API surface

sidecar answers

.flint/publish.ack
ok ⇒ the bytes are in S3 and
a sibling can read them

what it costs — MEASURED, loopback floors (0b); a latency-bound proxy moves these

bytes	checkout 3.3 s/GiB · publish 8.0 s/GiB
100k files	checkout 49.5 s · first publish 65 s (was 117 s before bounded fan-out) · idle tick 1.85 s
1M files	checkout 7 m 05 s · manifest 264 MiB at ~277 B/entry, exactly linear — the manifest is the wall
v1 cap	~250k files, and it falls out of the manifest size, not the byte count
idle costs ~5 bucket requests per tick, and the verbs add ZERO when unused — measured 20 vs 20 over 22 s	

worked example — a coding-agent fleet

A 2 GiB repo, 40k files, 30 agent pods, each on its own prefix. Checkout ≈ 7 s of bytes plus ~20 s of file count, so a pod is working inside half a minute. Each agent publishes when it finishes a task rather than on a timer, so a reviewer reads whole units of work, never a half-written tree. Idle pods cost about five requests a tick each. Nothing is shared live: if two agents must edit one tree simultaneously, that is the hub's job, not this one — and the lease will tell you so by refusing the second writer rather than merging them badly.

when to walk away from lean

- The tree does not fit the pod's disk, or is TiB-scale, or you need lazy access to a small slice of a huge dataset — FUSE.
- Two writers must see each other's bytes live, or you need byte-range locks, O_EXCL or shared sqlite/git ACROSS pods — the hub.
- A reader cannot tolerate seeing the last boundary instead of the last write — the hub. Gated mode narrows WHAT a reader sees, never HOW STALE it is.

The fit question is one question, and it is about the disk. Everything else follows: because the whole tree is local, the app gets real POSIX at local speed and the FUSE risk class disappears; because it is local, it must fit; and because publishing is a boundary rather than a write-through, a reader sees coherent points instead of a live tree. **The file count binds before the byte count.** At ~277 B per manifest entry the document is exactly linear, so a million small files costs a 264 MiB manifest while a million bytes costs nothing — which is why the v1 cap is expressed in files (~250k) and why the levers that mattered were fan-out and skipping the manifest parse on an idle tick, not raw throughput. **The verbs change the shape of the workflow, not its cost.** Cadence alone forces a choice between publishing too often (wasted work) and too rarely (a reviewer reads a half-finished tree); a declared boundary removes the choice, and costs nothing when unused. **Say no early.** The three questions above are ordered so the cheapest disqualifier comes first, and a workload that fails any of them is better served by another front end over the same bucket format — all three are mutually convertible, so choosing lean today does not strand the data if the answer changes.

The obvious question once several agents share a repo: how do they integrate? The answer is that they do not integrate through flint-lean at all. **Use plain git against a real git host, and let flint-lean keep the workspace durable.** Two mechanisms, two network paths, two different jobs — and they never touch, which is exactly why nothing special is needed to make them coexist. Full guide: docs/flint-lean-git-workflow.md.

